

# Fantastic Data Consistency Techniques and Where to Find Them

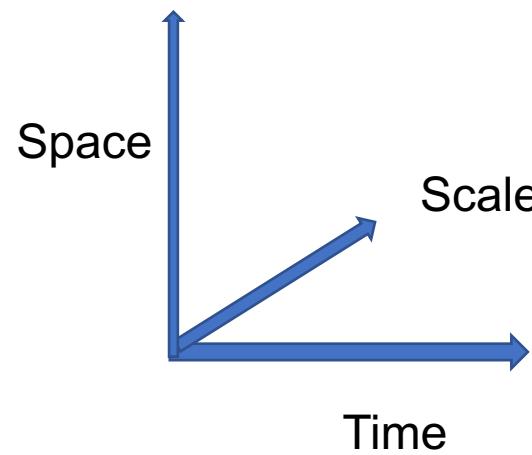
---

Gordon Hutchison  
IBM Development  
OpenLiberty  
@gordhut

September 2019

- This session covers data consistency available today for distributed microservices.
- Globally consistent data is no longer a safe assumption, and what will ultimately replace it is not yet resolved. Are the techniques that are becoming popular now such as eventual consistency, sagas, and CQRS/event sourcing frameworks here to stay? Or will they be displaced, in turn, by techniques emerging in multiple-editor shared document systems such as conflict-free replicated data types (CRDTs)?
- If you are wondering how to cope without the ACID properties of transactions in distributed data as you move to microservices. This session will inform you of the palette of solutions and libraries currently available.

# Consistency



## Local Scale

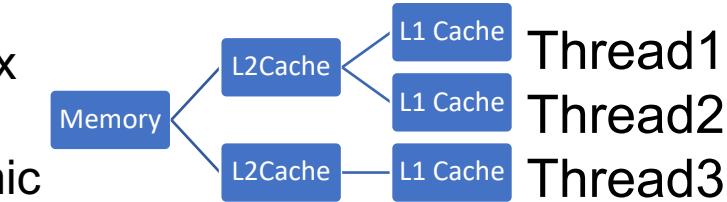
Cache-Memory

Weakly Consistency

POSIX mutex

Across Java Threads

Volatile/Atomic



Persistent Storage

Recoverable results

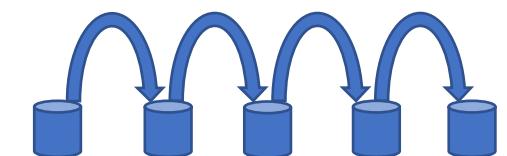
Repeatable Reads/Cursor Stability,

## Across Time

Always consistent view Databases -

A C ID

Eventually Consistent



## Across Space

Clusters – Consensus via Paxos etc,

XA – Coordinator+voting

Microservice Constellations



# Distributed Microservices

An application composed from co-operating services.

Services may be taking part in multiple applications.

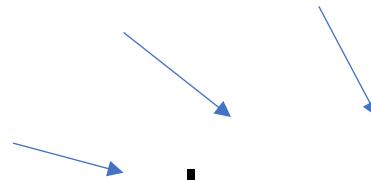
Services may be running:

- On-prem
- Cloud
- 3<sup>rd</sup> Party

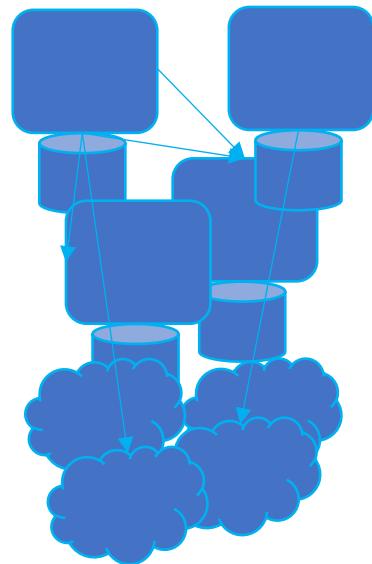
Also

- ‘Cattle not pets’
- Scale out rather than scale up
- Sharding
- Re-hashing
- Geo-replication

All this leads to



Increased Data Distribution



# Why look beyond distributed transactions?

Cross server call/data flow is:

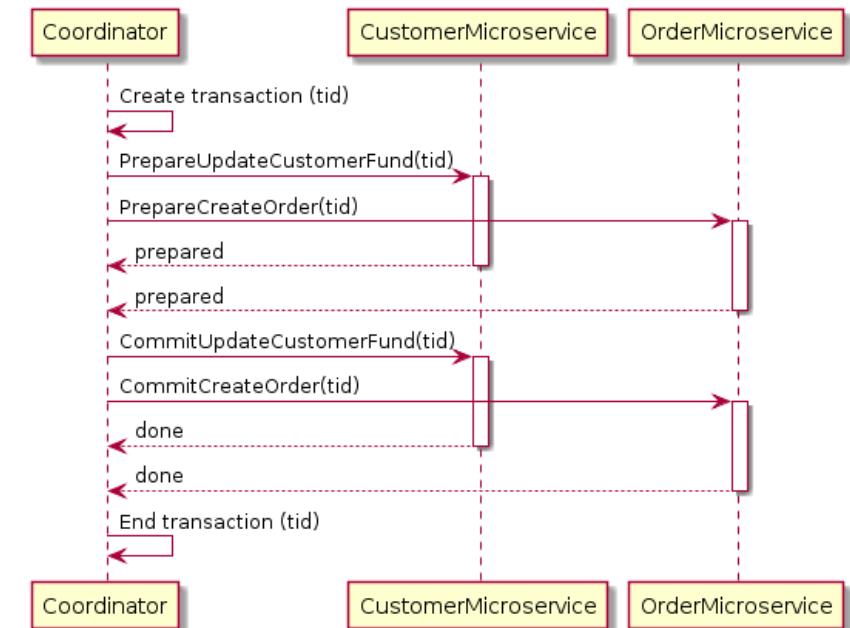
- finer grained
- more frequent
- more participants

Distributed transactions ‘hardening’  
overhead more costly.

Middleware:

- multi-vendor
- rapid turnover
- polyglot-technologies

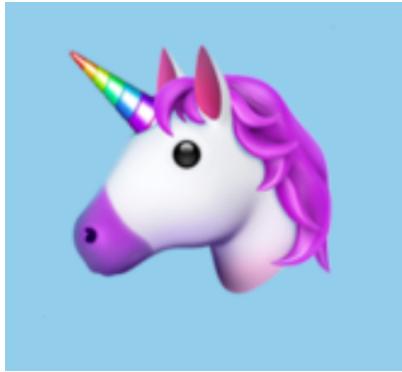
Distributed transactions/XA not supported  
by middleware



Blocking – Synchronous ‘long’ calls + locks...**blocking increases latency**

Design of data/lock ordering needed to avoid...**deadlock**

# Consistency for Microservices



Bandwidth  
'Cattle not pets'  
Polyglot software  
Data as an asset  
Consistency for 'free'



Latency  
Blocking  
Call outs  
Server affinity  
Software lock-in  
Data inconsistency

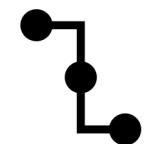
# Avoiding Contention Altogether?



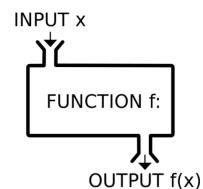
Optimistic Locking



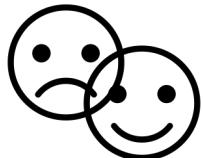
Sharding Aggregates with a Consistent Hash



Thread sharding and non-preemptable tasks



Functional Programming



Actors (Akka)

## Funny cartoon on Eventual Consistency

<https://twitter.com/mykola/status/1101337299525267457>

# Eventual Consistency

## CAP Theorem

What can services rely on?

What has actually happened locally.

Other parties informed (events)

Reconciliation is asynchronous to write

Weaker than

linearizable (read ALL completed writes)

strict serializability (of transactions)

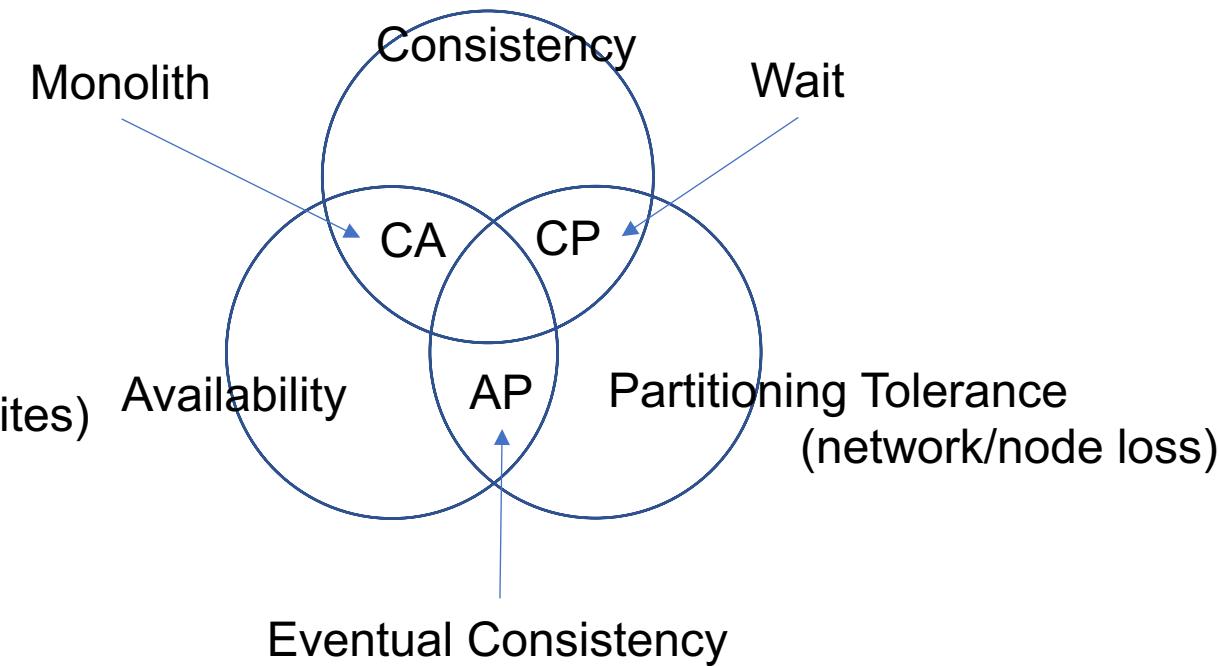
Commands (versus Queries)

Resolve conflicts by...

Consensus (coordinator)

Backing up (sagas undo/redo)

Distributed common rules (no 'conflicts')



# Long Running Activities and Compensation

Sagas.

- Server Local Transactional Steps
- Expose Intermediate State Tentatively
- Transactional outbox for event distribution
- Backtrack if necessary via compensations

Sub-step-**A**tomicity,

Local-**C**onsistency

Lumpy/leaky/temporary-**I**solation

**D**urability only after final commit



Microservices.io: saga (Eventuate)

Choreography - each local transaction publishes events that trigger other services

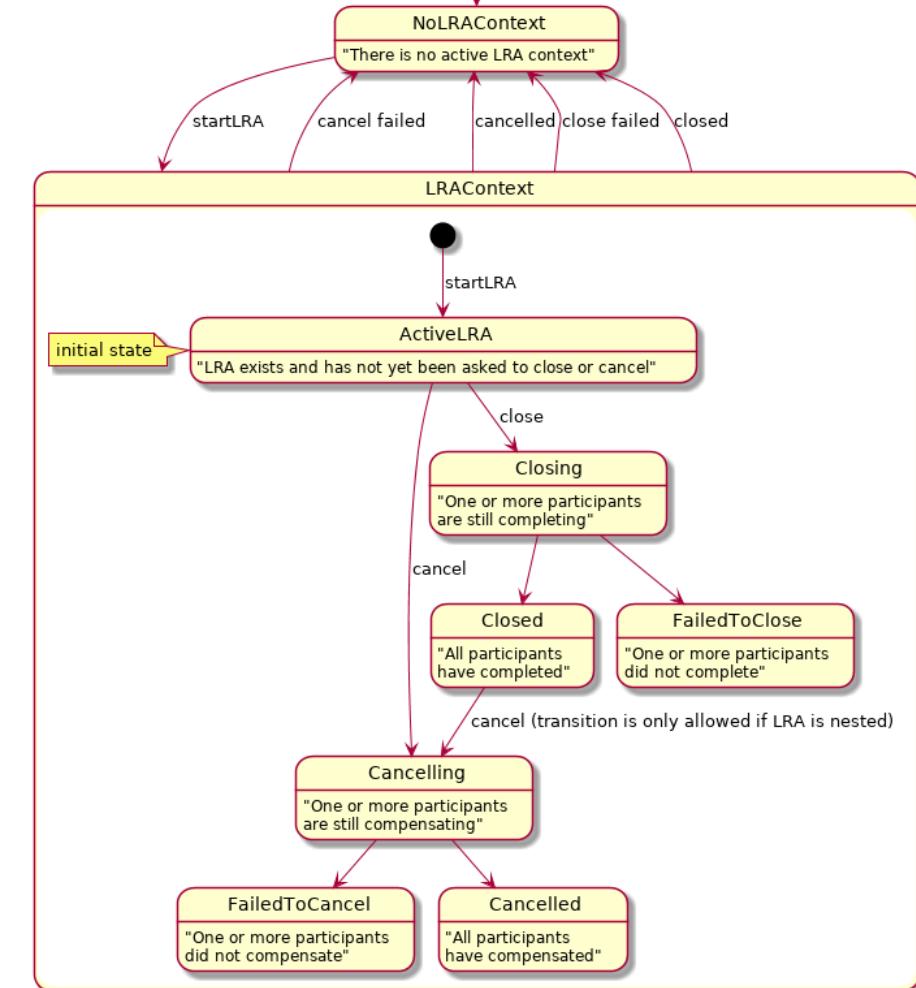
Orchestration - an orchestrator tells the participants what to execute

# Microprofile Long Running Activities. (LRA)

Annotation	Description
@LRA	Controls the life cycle of an LRA.
@Compensate	Indicates that the method should be invoked if the LRA is cancelled.
@Complete	Indicates that the method should be invoked if the LRA is closed.
@Forget	Indicates that the method may release any resources that were allocated for this LRA.
@Leave	Indicates that this class is no longer interested in this LRA.
@Status	When the annotated method is invoked it should report the status.

```

@Compensate @Path("/compensate")
@PUT public Response compensateWork(
    @HeaderParam(LRA_HTTP_CONTEXT_HEADER) URI lrald,
    String userData) {
    return Response.ok(
        ParticipantStatus.Compensated.name()).build();
}
  
```



# CQRS Architectures and Frameworks

Commands

- update state

Queries

- do not update state

Any microservice can be scaled independently

>1 view of data

Routing - Axon splits out

Commands(persistent-hash),  
Events(1->N),  
Queries(1ofN)

(Often associated with DDD and **Event Sourcing**)  
(handled by Aggregates)

Commands and Queries' state have different consistency needs.

Commands Update State

Single location – transactional

XA

or eventually consistent event sourced



Axon

# Event Sourcing

All (stimuli of) changes to state captured via an ‘event’

Events are immutable

Events are logged/stored  
indexed by time, consistency points

State of system can be arrived at solely by processing each event.

Events can be resent and rerouted  
(cf reactive manifesto)

Groups of events sent or not sent (‘transactional’)  
Event logs can be augmented by snapshots.



Kafka  
Apache Pulsar  
Axon Event Bus/Store

Eventuate  
EventStore  
DB2 Event Store  
(Cloud vendor equivalents)

# Change Data Capture (CDC)

Work from 'underneath' the database

Extracts changes to data from database logs/journals.

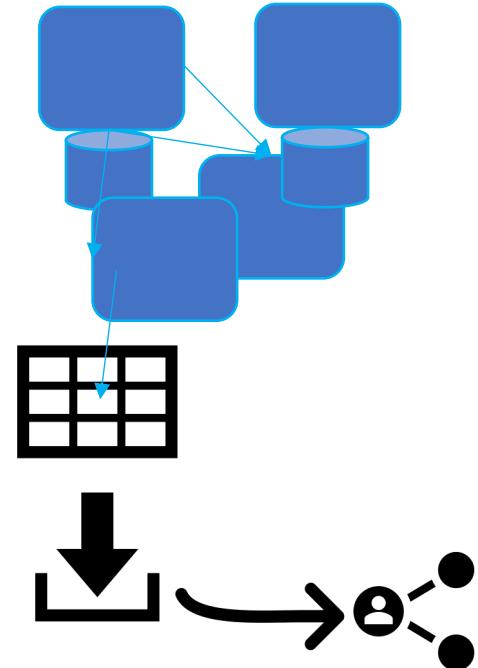
Whitelist-Publish → events → Subscribers

Change Events sent reliably over, for example, Kafka

Can be preceded by opening a Cursor over a table (timestamp )

Once cursor's data is transmitted

Change Events in log/journal from timestamp are forwarded



Database vendor solutions

( e.g. Oracle CDC/GoldenGate, IBM Infosphere, Azure )

Debezium



We've got the events sorted, how to we automate convergence?



# CALM Conjecture

**Consistency And Logical Monotonicity (CALM).**

*A program has an eventually consistent, coordination-free execution strategy if and only if it is expressible in monotonic datalog.*

<http://db.cs.berkeley.edu/papers/cidr11-bloom.pdf>

(Datalog is a Prolog like language for describing data operations

<https://en.wikipedia.org/wiki/Datalog> )



# Conflict-free Replicated Data Type

Fundamental Data Type:

built up from basic data structures:

flag, register, set, list, graph, tree, document, json

Forms of replication:

a local replica (work offline and resync later)

multiple editors (e.g. Boxnote, GoogleDoc )

multiple copies/devices ( clustered state )

Conflict Free:

'Collisions' still exist but...

A fixed set of rules exist for resolving any collisions

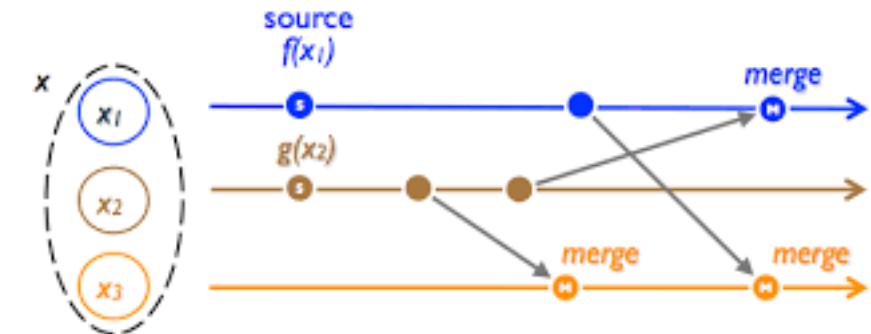
- without the need for consensus resolution.

e.g. Timestamped last writer wins

GPS/Atomic clocks, UUIDs

Convergence...towards a result

- not necessarily the exact result any 1 actor wants!



Domain level  
resolution still  
sometimes  
needed!

# CRDTs continued...

...convergence

Data structures progress ‘monotonically’ towards convergence:

**Operations** ‘concurrent’ ones are commutative  
(can arrive in different orders at different sites)

or...

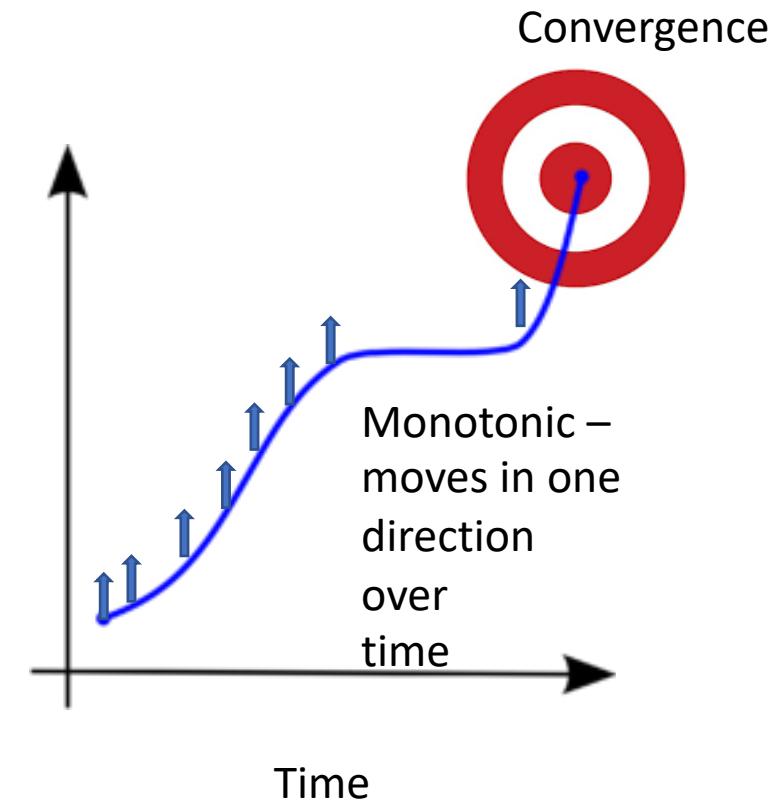
**State** forms a ‘semi-lattice’  
Internal data-structure to support this, indexing  
Retired UUIDs (Tombstones) etc.  
A ‘view’ function separates internal/external view

Causal Consistency

Born->Live->Die

Operations ‘pend’ until their prereq data/ops’ UUIDs turn up

Array [X,Y,Z] represented by [X<-Y<-Z] – “insert after” operations



# Interesting Commercial Libraries

(Not a personal or professional recommendation for any listing)

CDC

Debezium

CQRS

Axon

Eventuate

CRDTs

Lightbend

Akka distributed data

“Stateful Serverless”: CloudState

Amazon Dynamo/Apache Cassandra

Redis Enterprise - <https://redislabs.com/?s=CRDT>

Riak - <https://riak.com/distributed-data-types-riak-2-0/>

Microsoft Cosmos DB

“multiple, well-defined conflict resolution modes  
that guarantees strong convergence.”



Database!

# Let's talk open source, Java, Cloud and ... cool socks at booth #3101

- **Hands-on QuickLabs**
- Cloud-native development with Kabanero
- Getting started with Open Liberty
- Akka and Kubernetes: a symbiotic love story
- Intro to Kabanero and Spring

## Open Liberty

Experience the power of this cloud runtime for building cloud native apps in a fun, interactive game, Liberty Bikes. Open Liberty provides fully compatible implementations of Jakarta EE and Eclipse MicroProfile.

## Cloud-native Java

Talk to IBMers who are actively contributing to the open Java communities including: Open Liberty, Eclipse OpenJ9, Jakarta EE, Eclipse MicroProfile, Spring, Reactive and more.

## Journey to Cloud

Explore open source Kabanero which brings together Knative, Istio, and Tekton with new open projects Codewind, Appsody, and Razee into an end-to-end solution to architect, build, deploy, and manage the lifecycle of Kubernetes-based applications.